

ANUBHAV JAISWAL

+91 9899325122 +44 7304282326 anubhav.jaiswal99@hotmail.com [anubhav-jaiswal](#) [anubhavj99](#)

Technical Skills

- **Programming Langs:** Java (primary backend), Python (ML/data/automation), TypeScript/JavaScript(Node, React).
- **Cloud computing:** Lambda, ECS, DynamoDB, RDS, S3, CloudFront, AppConfig, Glue, Athena, ElastiCache, OpenSearch/Elasticsearch, CloudWatch, Route53, VPC, Docker. • **Frontend:** React, React-Native-Web, SSR/streaming HTML.
- **System Design:** Microservices & orchestration, plugin architectures, streaming pipelines, Caching strategies, Data Structures and Algorithms, fault tolerance, rate limiting & circuit breaking, observability & federated alarms, CI/CD.
- **Data/LLM:** SageMaker (training, fine-tuning **Mistral-7B**), Bedrock (**Claude**), LangChain, sentiment analysis pipelines.

Experience

Software Development Engineer II

Amazon — Rufus (Conversational AI Platform)

Nov 2024 – Present

London, United Kingdom

- **Orchestration & Scale (Backend first):** Re-architected Tier-1 Rufus service into a **plugin-based orchestration framework** with a **singleton orchestrator**, **hierarchical plugin state**, and **federated alarms** → enabled away-team onboarding in days (vs weeks), enabled away-teams to take full ownership of their codebase and business logic, and supported **35K+ TPS** in gamedays ensuring production readiness.
Tech: Core Java, AWS CloudWatch (with IaC), Simthy Models, architectural patterns (orchestration, plugins), internal load tools, load balancing.
- **Prod-like Dev Environments & CI/CD:** Established **shadow prod + developer pipelines**, cutting deploy time 3h → 15m (92% faster) and enabling **parallel dev stacks**, reduced **OE churn** and unblocked **rapid iteration**. Built one-click pipelines for quick dev environment setup.
Tech: CI/CD pipelines, internal release tooling, CloudWatch, ECS, Load Balancing, Route53, Docker
- **Reliability & Load Readiness:** Built end-to-end gameday infra routing; migrated legacy **load gen** → **repeatable high-TPS tests** validating production readiness at 35K+ TPS across tenants.
Tech: AWS, internal load tooling, observability/alarms
- **Streaming & Latency (UPL Tiger Team):** Optimized LLM streaming pipeline (chunk-wise processing, queueing, DOM batching) → >300ms per-section render reduction and 5–10s faster complete answer latency.
Tech: TypeScript/Node, streaming pipelines, perf profiling
- **Progressive Markdown Rendering (platformized):** Shipped SSR + progressive client patching for complex markdown (tables, nested lists, images) → project planned for 60 days, trimmed to 45, delivered in 7 days by prioritizing critical features, balancing speed (Bias for Action) with post-delivery backlog handling to maintain high system quality (Insist on Highest Standards); Tech reused by all other teams to render model data on Amazon Retail website.
Tech: TypeScript/React, SSR/HTML patching, streaming parsers
- **Progressive Reveal Algorithm (Patent in progress):** Enabled typing animation for HTML elements, supported 10× larger content while holding SLAs; robust partial→complete replacement logic. Implemented HTML patching with server side rendered HTML to reduce UPL.
Tech: TypeScript/React, DOM diffing, streaming state machines
- **Discovery UX — “Select & Ask Rufus” (innovation):** Prototyped contextual text-selection follow-ups to boost multi-turn engagement, Rufus discoverability and assistant-like UX.
Tech: TypeScript/React, prompt engineering
- **Leadership:** Mentored SDEs on plugin architecture & rendering pipelines; guided away teams (Deep Research, Papyrus, Customer Reviews, Returns) through service onboarding and delivery.
- **Recognition:** Top-Tier (Exceeds High Bar) in 2024–25 performance stack ranking.

Software Development Engineer II

Amazon — ShipTech

Apr 2022 – Nov 2024

Gurugram, India

- **Core Orchestration Owner:** Owned front-facing service to orchestrate transportation requests; enriched execution docs via 20+ downstreams using recursive dependency evaluation and dynamic config generation.
Tech: Java, AppConfig, DynamoDB, RDS, ElastiCache, CloudWatch
- **Cache Evaluation & Recursive Enrichment:** Framework to decide which dependencies to invoke from pre-cached state; recursively built the Transportation Document (TD) for correctness with minimal downstream load.
Tech: Java, DynamoDB, Glue, Athena
- **Employee Feedback Intelligence (LLM/ML):** Reduced leadership reporting 3 weeks → 15 mins (98 % faster) via sentiment analysis + actionable extraction and a chat interface; tracked action-item progress.
Tech: Python, SageMaker (fine-tuned Mistral-7B), OpenSearch, Lambda, CDN, CloudWatch

- **Human Query → DB Analysis Framework:** Designed inter-query decomposition and NL→SQL over massive datasets; approved by Amazon ML community and piloted on two databases (employee feedback & transportation). Auto-vended charts via S3 + CloudFront.
Tech: Python, LangChain, Bedrock (Claude), RDS/SQL, S3, CloudFront
- **Intake Process Automation:** Automated ShipTech intake across 50+ teams (team mapping, scope estimation, hallucination guardrails) → 70 %+ PM/PA effort reduction; integrated via Greasemonkey into SIM.
Tech: Python, SageMaker (Mistral-7B), Bedrock (Claude)
- **Contingency Management:** Ingested news/weather/geo streams to predict strikes/events and recommend actions; integrated with transport systems; continuous learning from ops feedback → improved SLA adherence.
Tech: Python, Lambda, OpenSearch, external data APIs
- **Reliability Guardrails:** Implemented rate-limiting guided by client-rate/ROI, circuit breakers for nth-level failures, and an issue-visibility tool (automated log aggregation) → faster triage, reduced on-call churn.
Tech: Java, CloudWatch, internal observability
- **Org-wide Influence:** Authored proposal for native multi-package shipment support across ShipTech/MMTT/LastMile → prioritized in 2024 OP.
- **Recognition:** Top-Tier (Exceeds High Bar) in 2023–24 performance stack ranking.

Software Development Engineer I

Amazon — ShipTech

Jul 2020 – Mar 2022

Gurugram, India

- Migrated transportation monolith → federated architecture, improving scalability and developer velocity.
Tech: Java, service decomposition, internal RPC
- Built external API interface for third-party shippers, translating global JSON → internal graph doc to enable external logistics integrations.
Tech: Java, API gateways, data transformation
- Implemented multi-package single-leg delivery capability, simplifying downstream execution.
Tech: Java, orchestration, validation

Software Development Engineer Intern

Amazon — MMPT

Jun 2019 – July 2019

Gurugram, India

- Migrated notification system to AWS-native: Lambda, ECS, DynamoDB, OpenSearch/Elasticsearch, ElastiCache, Route53, VPC → improved scalability, observability, and reduced infra overhead.

Education

Indraprastha Institute of Information Technology (IIIT), Delhi

Bachelor of Technology in Computer Science

2016 – 2020

8.47 GPA

Personal Projects

MVP for an automated human-like call | Bedrock / Claude, OpenAi / GPT, Twilio, ElevenLabs

Sept 2023

- Hosted a phone number at Twilio that you can either call or receive a call. Created 2 use-case including ticket booking for flights and Payment delay calls from Banks.
- Used OpenAi-Whisper model to transcribe offline to reduce latency for the response. Used GPT apis with 3.5-turbo model to get the response for the user. Used NLTK / Scikit libraries to determine the intent of the user response. In addition to the first system prompt to set the context, provided system roles in between human and assistant to guide GPT to provide better responses.

RDFLib QueryBuilder

Jan 2020

- Contributed to the OpenSource RDF library.
- Created a builder structure to create a query for RDF using Builder Pattern in code.
- Improved usage of the library and enhanced developer experience since now the queries could be understood by code formatters / IDEs and instead of a simple string.

Achievements and Extracurricular

- **Amazon Customer Obsession Award (Director-level org)** - Awarded for 2022 and 2023
- **Amazon Transportation Systems Hackathon** - Winner for 2022 and 2023
- **National Association for Blind (NAB)** - worked at an NGO for the blind in its Delhi branch. Taught children and helped them for chores that required compulsion of eyesight.
- **Play Tennis and Cricket as a hobby** - Have played junior state level for Delhi state